

How I ran my jupyter notebook

- I used NYU HPC Jupiter notebook from interactive apps to run my notebook with 14 cores and GPU because Google Codelab free version does not have enough memory space.

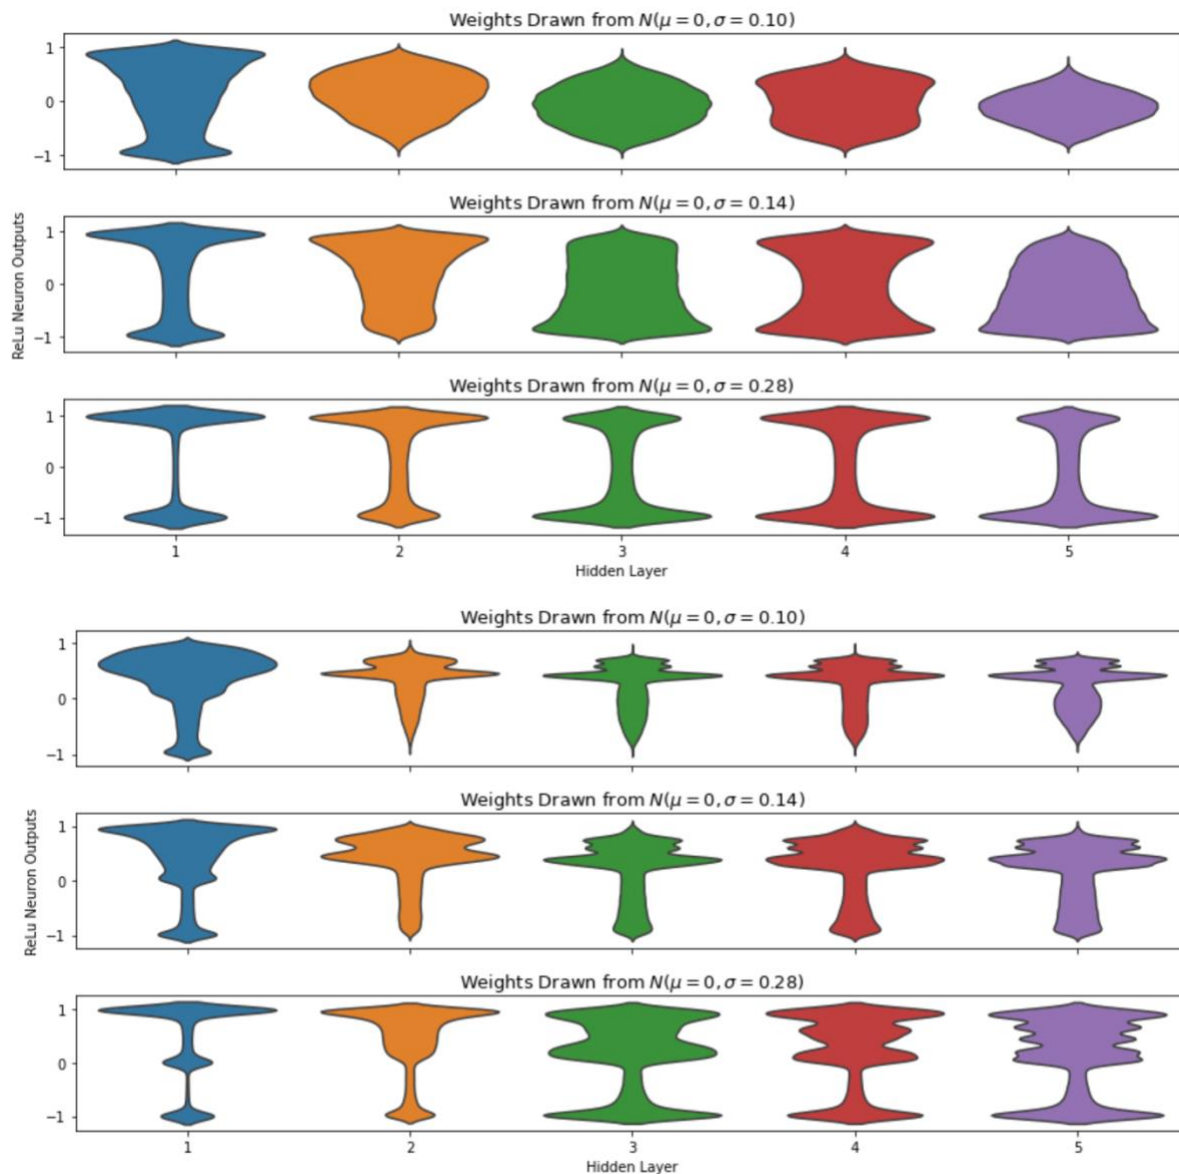
Google Codelab Link:

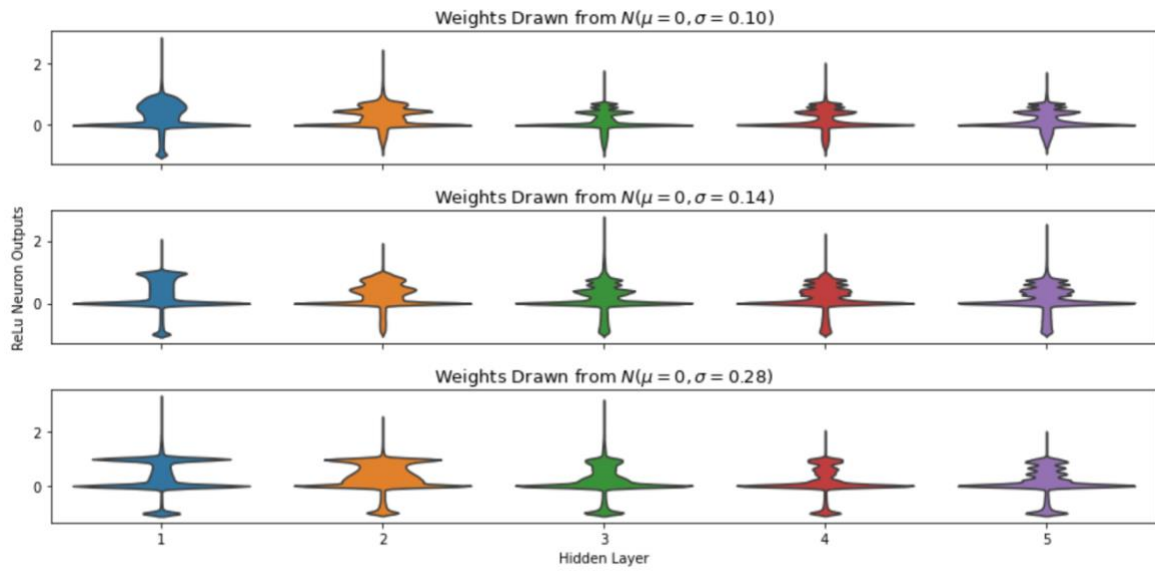
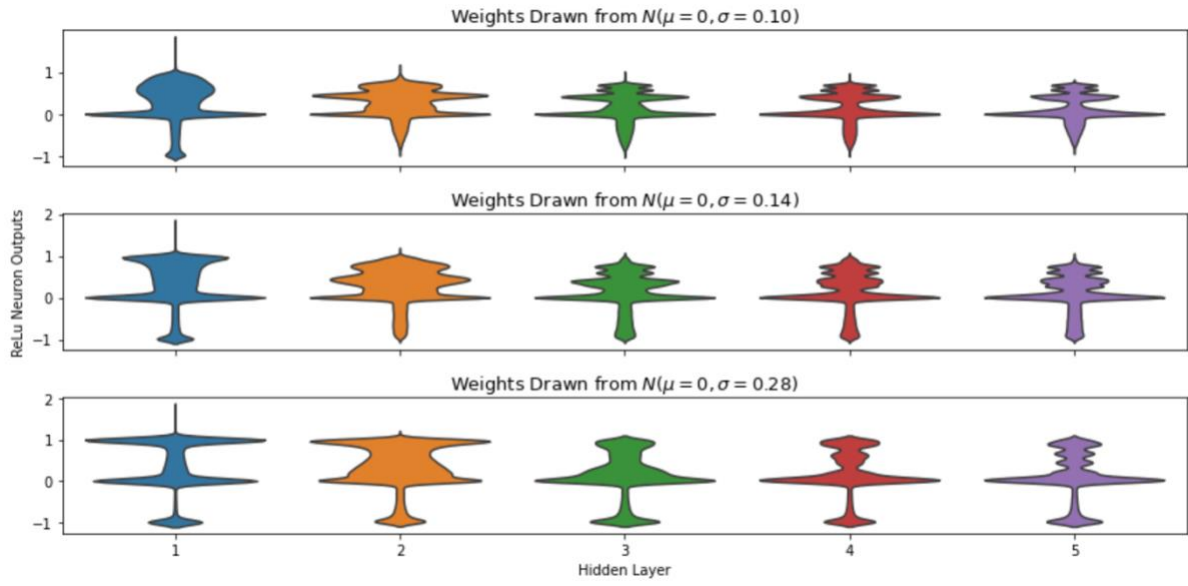
<https://colab.research.google.com/drive/1KZoqn3Xmqheo4ztORpWITsB79c5Zujjk?usp=sharing>

- Since I couldn't submit my Jupyter notebook on Brightspace, I uploaded all the files to Google drive.

- o <https://drive.google.com/drive/folders/1IYLcq5-b3Elr26pW5FNmU74AhTqFaqDy?usp=sharing>

Problem 1

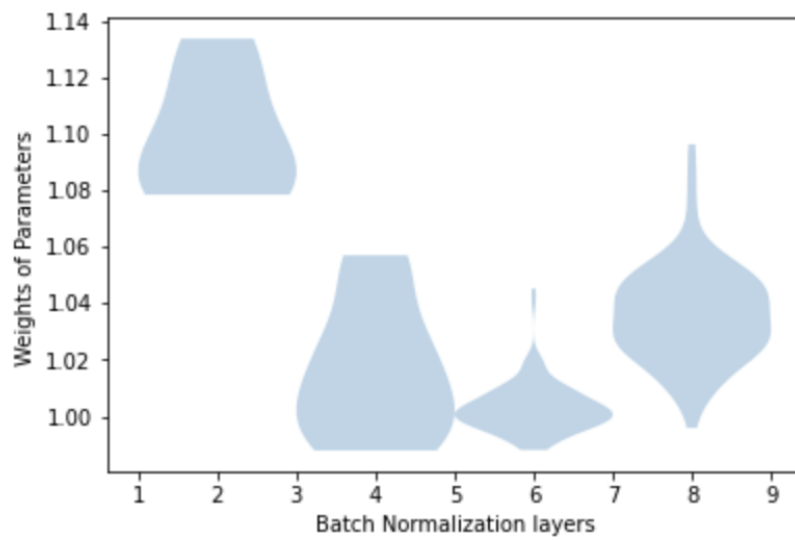




Problem 2

Part 2


- Plot the distribution of learned batch norm parameters for each layer



Problem 3

Part 1

- Plot the training loss as a function of the learning rate.




Part 2

- Plot train/validation loss and accuracy curve



Part 3




Problem 4

Part 2 without dropout

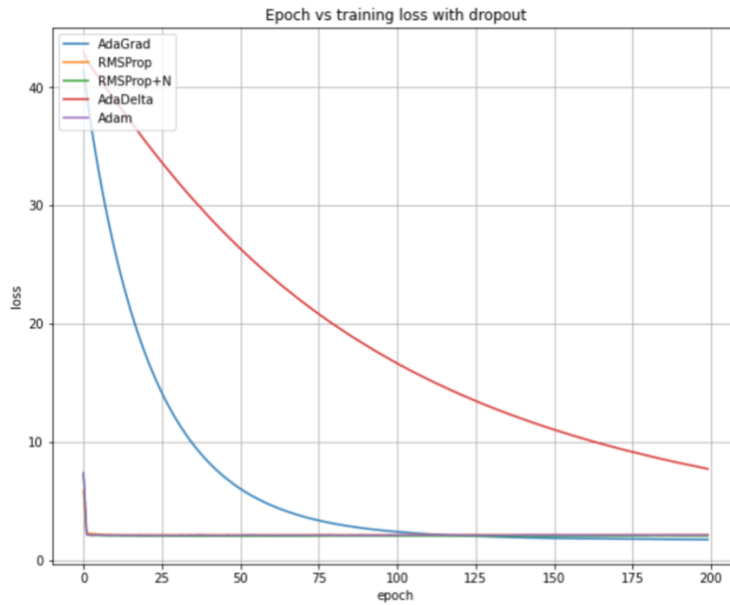
- plot the training loss vs the number of epochs.

lowest training loss
AdaGrad: 1.6244221925735474
RMSProp: 1.8377635478973389
RMSProp+N: 1.5503463745117188
AdaDelta: 7.079044342041016
Adam: 1.6213159561157227



Part 3 with dropout

lowest training loss
AdaGrad: 1.7369143962860107
RMSProp: 2.115575075149536
RMSProp+N: 2.028042793273926
AdaDelta: 7.716529846191406
Adam: 2.0822479724884033



Part 4

```

Test accuracy - models without dropout
313/313 [=====] - 1s 2ms/step - loss: 1.6799 - accuracy: 0.5031
Test Accuracy - Adagrad: 0.5030999779701233
313/313 [=====] - 1s 2ms/step - loss: 1.9086 - accuracy: 0.3608
Test Accuracy - RMSProp: 0.36079999804496765
313/313 [=====] - 1s 2ms/step - loss: 1.5688 - accuracy: 0.4755
Test Accuracy - RMSProp+Nesterov: 0.4754999876022339
313/313 [=====] - 1s 2ms/step - loss: 7.1394 - accuracy: 0.4537
Test Accuracy - Adadelta: 0.4537000060081482
313/313 [=====] - 1s 2ms/step - loss: 1.7330 - accuracy: 0.4200
Test Accuracy - Adam: 0.41999998688697815
Test accuracy - models with dropout
313/313 [=====] - 1s 2ms/step - loss: 1.6934 - accuracy: 0.4811
Test Accuracy - Adagrad: 0.4810999929904938
313/313 [=====] - 1s 2ms/step - loss: 2.0576 - accuracy: 0.2864
Test Accuracy - RMSProp: 0.2863999903202057
313/313 [=====] - 1s 2ms/step - loss: 2.0737 - accuracy: 0.2396
Test Accuracy - RMSProp+Nesterov: 0.23960000276565552
313/313 [=====] - 1s 2ms/step - loss: 7.6006 - accuracy: 0.3985
Test Accuracy - Adadelta: 0.398499995470047
313/313 [=====] - 1s 2ms/step - loss: 2.1113 - accuracy: 0.2138
Test Accuracy - Adam: 0.21379999816417694

```

Problem 5

Part 1

Layer	Number of Activations (Memory)	Parameters (Compute)
Input	$224*224*3=150K$	0
CONV3-64	$224*224*64=3.2M$	$(3*3*3)*64 = 1,728$
CONV3-64	$224*224*64=3.2M$	$(3*3*64)*64 = 36,864$
POOL2	$112*112*64=800K$	0
CONV3-128	$112*112*128=1605632$	$3*3*64*128=73728$
CONV3-128	$112*112*128=1605632$	$3*3*128*256=294912$
POOL2	$56*56*128=400K$	0
CONV3-256	$56*56*256=802816$	$3*3*128*256=294912$
CONV3-256	$56*56*256=800K$	$(3*3*256)*256 = 589,824$
CONV3-256	$56*56*256=802816$	$3*3*256*256=589824$
CONV3-256	$56*56*256=802816$	$3*3*256*256=589824$
POOL2	$28*28*256$	0
CONV3-512	$28*28*512=400K$	$(3*3*256)*512 = 1,179,648$
CONV3-512	$28*28*512=401408$	$3*3*256*512=1179648$
CONV3-512	$28*28*512=400K$	$3*3*256*512=1179648$
CONV3-512	$28*28*512=401408$	$3*3*512*512=2359296$
POOL2	$14*14*512=100353$	0
CONV3-512	$14*14*512=100353$	$3*3*512*512=2359296$
CONV3-512	$14*14*512=100353$	$3*3*512*512=2359296$
CONV3-512	$14*14*512=100353$	$3*3*512*512=2359296$
CONV3-512	$14*14*512=100353$	$3*3*512*512=2359296$
POOL2	$7*7*512=25088$	0
FC	4096	$4096*25088=102760488$
FC	4096	$4096*4096 = 16,777,216$
FC	1000	$4096*1000=4096000$
TOTAL	16.5M	140M

Table 1: VGG19 memory and weights

